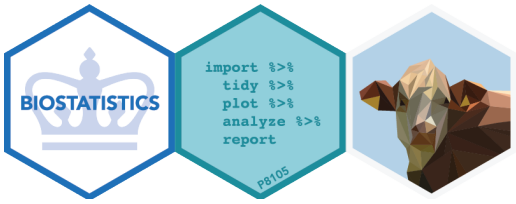

P8105: Data Science I

COURSE DESCRIPTION

Contemporary biostatistics and data analysis depends on the **mastery of tools for computation, visualization, dissemination, and reproducibility** in addition to proficiency in traditional statistical techniques. The goal of this course is to provide training in the elements of a complete pipeline for data analysis. It is targeted to MS, MPH, and PhD students with some data analysis experience.

LEARNING OBJECTIVES

Students who successfully complete this course will:

- Integrate the principles of data organization into their analyses;
- Easily produce static and interactive graphics;
- Implement analyses in a reproducible way;
- Use GitHub to publish and disseminate analyses;
- Develop usable software packages in R;
- Collect data from online sources using web-scraping.

INSTRUCTOR

Jeff Goldsmith, PhD
Associate Professor of Biostatistics
Email: <ajg2202@cumc.columbia.edu>

TEACHING ASSISTANTS

Angel Garcia de la Garza (co-lead)
Rui Huang
Christian Pasqual
Dana Zeng
Zelos Zhu
Erin McDonnell (co-lead)
Courtney Johnson
Jiayi Shen
Eleanor Zhang
Coco Zou

CLASS SESSIONS

Tuesdays and Thursdays 10:00 - 11:20, Physicians & Surgeons, Amphitheater 7
Two exceptions: lecture will be in Alumni Auditorium on 10/3 Hammer 401 on 10/22

DISCUSSION BOARD, OFFICE HOURS, EMAIL

There are several ways to get help answering course related questions. Slack channels will be created for each topic, and are a way to ask and answer questions in real time during class sessions. The course website through courseworks includes a discussion board, and we encourage students to proactively use this as a way to get help and to help others. For more complex issues, in-person office hours may be more appropriate (M-Th 11:30-12:50 in ARB 627; M-Th 4:30-5:50 in ARB 627; F 11:30-12:50 in ARB 657; and by appointment). Email should be used to address questions regarding course structure or policy; content-related questions will generally be referred to the discussion board or office hours.

PREREQUISITES
Experience in R programming (or programming in another language) and data analysis is **recommended but not required**. A laptop with R installed is required and should be brought to every class session.

RECOMMENDED REFERENCES (note: there are no required texts for this course)
The Internet (stackoverflow; google; blog posts; twitter)
R for Data Science by G. Grolemund and H. Wickham
Exploratory Data Analysis with R by R Peng
R Programming for Data Science by R Peng.
R Packages by H. Wickham
Advanced R by H. Wickham

ASSESSMENT AND GRADING POLICY
Student grades will be based on:
Homework Assignments 30%
Midterm Project.....35%
Final Project 35%
Questions regarding the grading of HW assignments must be raised within a week of the assignment being returned.

Homework assignments will be due following the completion of each course topic. Only electronic submissions will be accepted. Collaboration on homework assignments is governed by the course policy on collaboration. Late homework will not be accepted. Unclear or disorganized homework will have points removed, even if the content is correct.

The midterm project will focus on demonstrating proficiency in the topics covered in the first half of the course (R, R Markdown, data wrangling, exploratory analysis, and plotting). Collaboration on the midterm project is **strictly prohibited**.

The final project will consist of a complete analytic pipeline, starting with getting data and ending with a polished report, website, and screencast. This will be a group project, and group members will collaborate on the project using GitHub.

SOFTWARE USE
We will use R and R Markdown; R Studio is recommended.

COURES WEBSITE
The course website contains lecture materials, homework assignments, supplementary materials, helpful links, and project information. It can be accessed at www.p8105.com.

COURSE STRUCTURE
Class sessions will be lectures, delivered using a mix of static content and live demonstrations.

COURSE SCHEDULE

Lecture 1: What is data science?

Learning Objectives:

- Define "data science" and its role in public health research

Required Reading:

- "50 Years of Data Science" by David Donoho
- The Data Science Venn Diagram
- 'Janitor Work' vs 'Data Carpentry'
- 'What have you tried?' and a follow-up by the author
- *R Programming for Data Science:*
  - History and Overview of R
  - Getting Started with R

Homework: Assignment 0 (for details on all assignments, see below)

Lecture 2: Best Practices

Learning Objectives:

- Use best practices for coding, including commenting and human-readable naming structures.

Required Reading:

- *R for Data Science:*
  - 4) Workflow: basics, scripts, projects
- R Studio Code Diagnostics
- BEH Commandments for Variable Names
- Using R Projects

Homework: Assignment 1

Lecture 3: Writing with data

Learning Objectives:

- Implement basic analyses using R Markdown and R Notebooks. Export analysis reports into several formats.

Required Reading:

- *R for Data Science:*
  - 27.1 – 27.4) R Markdown
  - 29.1 – 29.5) R Markdown Formats
  - 30) R Markdown Workflow

Homework: Assignment 1

Lecture 4: Version control and dissemination

Learning Objectives:

- Create local and remote Git repositories, and integrate with R Projects. Use commits for version control.

Required Reading:

- Happy Git and GitHub for the useR

Homework: Assignment 1

Lecture 5: Data import

Learning Objectives:

- Read data into R from a variety of sources
- Parse variable types

Required Reading:

- *R Programming for Data Science:*
  - Getting Data In and Out of R
- *R for Data Science:*
  - 11) Data Import

Homework: Assignment 2

Lecture 6: Data manipulation

Learning Objectives:

- Clean and organize data using dplyr verbs and piping.

Required Reading:

- *R Programming for Data Science:*
  - Managing Data Frames with the dplyr package
- *R for Data Science:*
  - 12.1 – 12.5) Tidy Data
  - 18) Pipes

Homework: Assignment 2

Lecture 7: Tidy data and relational datasets

Learning Objectives:

- Explain principles of "tidy" data.
- Use relational databases; merging datasets

Required Reading:

- *R Programming for Data Science:*
  - Getting Data In and Out of R
  - Managing Data Frames with the dplyr package
- *R for Data Science:*
  - 11) Data Import
  - 12.1 – 12.5) Tidy Data
  - 18) Pipes

Homework: Assignment 2

Lecture 8: Data visualization

Learning Objectives:

- Create effective graphics using ggplot using the grammar of graphics. Implement best practices for effective graphical communication.

Required Reading:

- "A Layered Grammar of Graphics" by Hadley Wickham
- *R for Data Science:*
  - 3) Data Visualization
  - 28) Graphics for Communication

Homework: Assignment 3

Lecture 9: Data visualization

Learning Objectives:

- Create effective graphics using ggplot using the grammar of graphics. Implement best practices for effective graphical communication.

Required Reading:

- "A Layered Grammar of Graphics" by Hadley Wickham
- *R for Data Science:*
  - 3) Data Visualization
  - 28) Graphics for Communication

Homework: Assignment 3

Lecture 10: Exploratory analysis

Learning Objectives:

- Conduct exploratory analyses using dplyr verbs (group_by and summarize).

Required Reading:

- *R for Data Science:*
  - 7) Exploratory analysis

Homework: Assignment 3

Lecture 11: Case study

Learning Objectives:

- Pull together skills learned through this point
- Produce a complete analysis and written summary

Lecture 12: Reading Data from the Web

Learning Objectives:

- Gather data from online sources (i.e. "scrape") using APIs, rvest and httr.

Lecture 13: Strings and factors

Learning Objectives:

- Edit / manipulate strings; take control of factors

Lecture 14: Websites

Learning Objectives:

- Publish a personal website using GitHub Pages.

Required Reading:

- GitHub Pages

Homework: Assignment 4

Lecture 15: Plot.ly and dashboards

Learning Objectives:

- Create interactive graphics using plot.ly
- Design an data dashboard using flexdashboard

Homework: Assignment 4

Lecture 16: Writing R functions

Learning Objectives:

- Create simple R functions to abstract common processes.

Required Reading:

- *R Programming for Data Science:*
  - Functions
  - Scoping Rules of R
- *R for Data Science:*
  - 19) Functions

Homework: Assignment 5

Lecture 17: Simulating data

Learning Objectives:

- Simulate datasets in R. Use loops, apply functions, and map functions.

Required Reading:

- *R Programming for Data Science:*
  - Simulation
  - Loop functions

Homework: Assignment 5

Lecture 18: Simulation

Learning Objectives:

- Use loop and apply functions to simulate data. Explore statistical properties of usual estimate methods using simulations.

Required Reading:

- *R Programming for Data Science:*
  - Simulation
  - Loop functions

Homework: Assignment 5

Lecture 19: Linear and generalized linear models

Learning Objectives:

- Review fundamentals of linear and generalized linear models. Fit models in R and tidy results for further analysis.

Required Reading:

- *Introduction to Statistical Learning with R*
  - Chapter 3.1-3.3
  - Chapter 4.1.-4.3

Homework: Assignment 6

Lecture 20: Cross validation

Learning Objectives:

- Use cross validation to assess predictive value of a model. Implement CV using tools for iteration.

Required Reading:

- *Introduction to Statistical Learning with R*
  - Chapter 5.1

Homework: Assignment 6

Lecture 21: Bootstrapping

Learning Objectives:

- Implement the bootstrap to obtain inference in non-standard cases using tools for iteration.

Required Reading:

- *Introduction to Statistical Learning with R*
  - Chapter 5.2

Homework: Assignment 6

Lecture 22: Extra topics

Learning Objectives:

Required Reading:

Homework: Assignment 7

Lecture 23: Extra topics

Learning Objectives:

Required Reading:

Homework: Assignment 7

ASSIGNMENTS

Assignment 0	
L1	Assignment 0 covers the installation of software and creation of accounts.

Assignment 1	
L2-L4	Assignment 1 covers basic R coding, including variable assignments, data manipulation, and the use of basic functions. Submissions will use the R Markdown format to ensure reproducibility, and best practices for clarity.

Assignment 2	
L5-L7	Assignment 2 covers data input and output; principles of data cleaning; and implementation of data cleaning using dplyr.

Assignment 3	
L8-L10	Assignment 3 covers exploratory data analysis. Students are expected to produce reasonable summaries of data, including both tables and graphics, and accompany these with clearly-written text describing the results.

Assignment 4	
L14-L15	Assignment 4 covers dashboards and websites. Students will develop a professional website including their contact information and highlighting their work / CV. This website will also link to a dashboard.

Assignment 5	
L16-L18	Assignment 5 covers simulation and looping. Students will conduct simulation experiments to explore basic statistical properties, and will illustrate these graphically and in words.

Assignment 6	
L19-L21	Assignment 6 covers linear models.

Assignment 7	
L22-L23	Assignment 7 covers extra topics.

MAILMAN SCHOOL POLICIES AND EXPECTATIONS

Students and faculty have a shared commitment to the School's mission, values and oath. http://mailman.columbia.edu/about-us/school-mission/

Academic Integrity

Students are required to adhere to the Mailman School Honor Code, available online at http://mailman.columbia.edu/honorcode.

Disability Access

In order to receive disability-related academic accommodations, students must first be registered with the Office of Disability Services (ODS). Students who have, or think they may have a disability are invited to contact ODS for a confidential discussion at 212.854.2388 (V) 212.854.2378 (TTY), or by email at disability@columbia.edu. If you have already registered with ODS, please speak to your instructor to ensure that s/he has been notified of your recommended accommodations by Lillian Morales (lm31@columbia.edu), the School's liaison to the Office of Disability Services.

Student Affairs

The Office of Student Affairs (OSA) supports the needs of students who experience life challenges, which may disrupt their successful completion of a Public Health degree. Students' needs may manifest in such areas as their physical, mental, and/or emotional health; economic, family, and/or social stressors; difficulties resulting from adjustment to graduate-level work and/or transitioning to academia after time away from school; as well as other barriers to students' success. Students in need of support should reach out to OSA by phone (212-342-3128), email, or as a walk-in during office hours (8:00 a.m. – 6:00 p.m.; located on the 10th floor of ARB). Students may also directly access the resources and services of Student Health Services, Mental Health, Services, the Center for Student Wellness, and other supportive offices throughout CUMC directly through the offices' websites, links to which can be found on the Health and Wellness page of the Mailman website.